CHAPTER 19

Network Layer: Logical Addressing

As we discussed in Chapter 2, communication at the network layer is host-to-host (computer-to-computer); a computer somewhere in the world needs to communicate with another computer somewhere else in the world. Usually, computers communicate through the Internet. The packet transmitted by the sending computer may pass through several LANs or WANs before reaching the destination computer.

For this level of communication, we need a global addressing scheme; we called this logical addressing in Chapter 2. Today, we use the term **IP address** to mean a logical address in the network layer of the TCP/IP protocol suite.

The Internet addresses are 32 bits in length; this gives us a maximum of 2^{32} addresses. These addresses are referred to as IPv4 (IP version 4) addresses or simply IP addresses if there is no confusion.

The need for more addresses, in addition to other concerns about the IP layer, motivated a new design of the IP layer called the new generation of IP or IPv6 (IP version 6). In this version, the Internet uses 128-bit addresses that give much greater flexibility in address allocation. These addresses are referred to as IPv6 (IP version 6) addresses.

In this chapter, we first discuss IPv4 addresses, which are currently being used in the Internet. We then discuss the IPv6 addresses, which may become dominant in the future.

19.1   IPv4 ADDRESSES

An **IPv4 address** is a 32-bit address that *uniquely* and *universally* defines the connection of a device (for example, a computer or a router) to the Internet.

An IPv4 address is 32 bits long.

IPv4 addresses are unique. They are unique in the sense that each address defines one, and only one, connection to the Internet. Two devices on the Internet can never have the same address at the same time. We will see later that, by using some strategies, an address may be assigned to a device for a time period and then taken away and assigned to another device.

On the other hand, if a device operating at the network layer has *m* connections to the Internet, it needs to have *m* addresses. We will see later that a router is such a device.

The IPv4 addresses are universal in the sense that the addressing system must be accepted by any host that wants to be connected to the Internet.

> **The IPv4 addresses are unique and universal.**

Address Space

A protocol such as IPv4 that defines addresses has an **address space.** An address space is the total number of addresses used by the protocol. If a protocol uses *N* bits to define an address, the address space is 2^N because each bit can have two different values (0 or 1) and *N* bits can have 2^N values.

IPv4 uses 32-bit addresses, which means that the address space is 2^{32} or 4,294,967,296 (more than 4 billion). This means that, theoretically, if there were no restrictions, more than 4 billion devices could be connected to the Internet. We will see shortly that the actual number is much less because of the restrictions imposed on the addresses.

> **The address space of IPv4 is 2^{32} or 4,294,967,296.**

Notations

There are two prevalent notations to show an IPv4 address: **binary notation** and **dotted-decimal notation.**

Binary Notation

In binary notation, the IPv4 address is displayed as 32 bits. Each octet is often referred to as a byte. So it is common to hear an IPv4 address referred to as a 32-bit address or a 4-byte address. The following is an example of an IPv4 address in binary notation:

> 01110101 10010101 00011101 00000010

Dotted-Decimal Notation

To make the IPv4 address more compact and easier to read, Internet addresses are usually written in decimal form with a decimal point (dot) separating the bytes. The following is the **dotted-decimal notation** of the above address:

> 117.149.29.2

Figure 19.1 shows an IPv4 address in both binary and dotted-decimal notation. Note that because each byte (octet) is 8 bits, each number in dotted-decimal notation is a value ranging from 0 to 255.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

SECTION 19.1   IPv4 ADDRESSES   551

Figure 19.1   *Dotted-decimal notation and binary notation for an IPv4 address*

| 10000000 | 00001011 | 00000011 | 00011111 |

128.11.3.31

Numbering systems are reviewed in Appendix B.

Example 19.1

Change the following IPv4 addresses from binary notation to dotted-decimal notation.

   a. 10000001 00001011   00001011 11101111
   b. 11000001 10000011   00011011 11111111

Solution

We replace each group of 8 bits with its equivalent decimal number (see Appendix B) and add dots for separation.

   a. 129.11.11.239
   b. 193.131.27.255

Example 19.2

Change the following IPv4 addresses from dotted-decimal notation to binary notation.

   a. 111.56.45.78
   b. 221.34.7.82

Solution

We replace each decimal number with its binary equivalent (see Appendix B).

   a. 01101111 00111000 00101101 01001110
   b. 11011101 00100010 00000111 01010010

Example 19.3

Find the error, if any, in the following IPv4 addresses.

   a. 111.56.045.78
   b. 221.34.7.8.20
   c. 75.45.301.14
   d. 11100010.23.14.67

Solution

   a. There must be no leading zero (045).
   b. There can be no more than four numbers in an IPv4 address.
   c. Each number needs to be less than or equal to 255 (301 is outside this range).
   d. A mixture of binary notation and dotted-decimal notation is not allowed.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Classful Addressing

IPv4 addressing, at its inception, used the concept of classes. This architecture is called **classful addressing.** Although this scheme is becoming obsolete, we briefly discuss it here to show the rationale behind classless addressing.

In classful addressing, the address space is divided into five classes: A, B, C, D, and E. Each class occupies some part of the address space.

> **In classful addressing, the address space is divided into five classes:**
> **A, B, C, D, and E.**

We can find the class of an address when given the address in binary notation or dotted-decimal notation. If the address is given in binary notation, the first few bits can immediately tell us the class of the address. If the address is given in decimal-dotted notation, the first byte defines the class. Both methods are shown in Figure 19.2.

Figure 19.2   *Finding the classes in binary and dotted-decimal notation*

	First byte	Second byte	Third byte	Fourth byte
Class A	0			
Class B	10			
Class C	110			
Class D	1110			
Class E	1111			

a. Binary notation

	First byte	Second byte	Third byte	Fourth byte
Class A	0–127			
Class B	128–191			
Class C	192–223			
Class D	224–239			
Class E	240–255			

b. Dotted-decimal notation

Example 19.4

Find the class of each address.

   a. **0**0000001 00001011   00001011 11101111
   b. **110**00001 10000011   00011011 11111111
   c. **14**.23.120.8
   d. **252**.5.15.111

Solution

   a. The first bit is 0. This is a **class A address.**
   b. The first 2 bits are 1; the third bit is 0. This is a **class C address.**
   c. The first byte is 14 (between 0 and 127); the class is A.
   d. The first byte is 252 (between 240 and 255); the class is E.

Classes and Blocks

One problem with classful addressing is that each class is divided into a fixed number of blocks with each block having a fixed size as shown in Table 19.1.

Table 19.1 *Number of blocks and block size in classful IPv4 addressing*

Class	Number of Blocks	Block Size	Application
A	128	16,777,216	Unicast
B	16,384	65,536	Unicast
C	2,097,152	256	Unicast
D	1	268,435,456	Multicast
E	1	268,435,456	Reserved

Let us examine the table. Previously, when an organization requested a block of addresses, it was granted one in class A, B, or C. **Class A addresses** were designed for large organizations with a large number of attached hosts or routers. **Class B addresses** were designed for midsize organizations with tens of thousands of attached hosts or routers. **Class C addresses** were designed for small organizations with a small number of attached hosts or routers.

We can see the flaw in this design. A block in class A address is too large for almost any organization. This means most of the addresses in class A were wasted and were not used. A block in class B is also very large, probably too large for many of the organizations that received a class B block. A block in class C is probably too small for many organizations. **Class D addresses** were designed for multicasting as we will see in a later chapter. Each address in this class is used to define one group of hosts on the Internet. The Internet authorities wrongly predicted a need for 268,435,456 groups. This never happened and many addresses were wasted here too. And lastly, the **class E addresses** were reserved for future use; only a few were used, resulting in another waste of addresses.

> In classful addressing, a large part of the available addresses were wasted.

Netid and Hostid

In classful addressing, an IP address in class A, B, or C is divided into **netid** and **hostid.** These parts are of varying lengths, depending on the class of the address. Figure 19.2 shows some netid and hostid bytes. The netid is in color, the hostid is in white. Note that the concept does not apply to classes D and E.

In class A, one byte defines the netid and three bytes define the hostid. In class B, two bytes define the netid and two bytes define the hostid. In class C, three bytes define the netid and one byte defines the hostid.

Mask

Although the length of the netid and hostid (in bits) is predetermined in classful addressing, we can also use a **mask** (also called the **default mask**), a 32-bit number made of

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

contiguous 1s followed by contiguous 0s. The masks for classes A, B, and C are shown
in Table 19.2. The concept does not apply to classes D and E.

Table 19.2 *Default masks for classful addressing*

Class	Binary	Dotted-Decimal	CIDR
A	**11111111** 00000000 00000000 00000000	**255**.0.0.0	/8
B	**11111111 11111111** 00000000 00000000	**255.255**.0.0	/16
C	**11111111 11111111 11111111** 00000000	**255.255.255**.0	/24

The mask can help us to find the netid and the hostid. For example, the mask for a
class A address has eight 1s, which means the first 8 bits of any address in class A
define the netid; the next 24 bits define the hostid.

The last column of Table 19.2 shows the mask in the form */n* where *n* can be 8, 16,
or 24 in classful addressing. This notation is also called slash notation or **Classless
Interdomain Routing (CIDR)** notation. The notation is used in classless addressing,
which we will discuss later. We introduce it here because it can also be applied to class-
ful addressing. We will show later that classful addressing is a special case of classless
addressing.

Subnetting

During the era of classful addressing, **subnetting** was introduced. If an organization was
granted a large block in class A or B, it could divide the addresses into several contiguous
groups and assign each group to smaller networks (called **subnets**) or, in rare cases,
share part of the addresses with neighbors. Subnetting increases the number of 1s in the
mask, as we will see later when we discuss classless addressing.

Supernetting

The time came when most of the class A and class B addresses were depleted; however,
there was still a huge demand for midsize blocks. The size of a class C block with a
maximum number of 256 addresses did not satisfy the needs of most organizations.
Even a midsize organization needed more addresses. One solution was **supernetting.**
In supernetting, an organization can combine several class C blocks to create a larger
range of addresses. In other words, several networks are combined to create a super-
network or a **supernet.** An organization can apply for a set of class C blocks instead of
just one. For example, an organization that needs 1000 addresses can be granted four
contiguous class C blocks. The organization can then use these addresses to create one
supernetwork. Supernetting decreases the number of 1s in the mask. For example, if an
organization is given four class C addresses, the mask changes from /24 to /22. We will
see that classless addressing eliminated the need for supernetting.

Address Depletion

The flaws in classful addressing scheme combined with the fast growth of the Internet led
to the near depletion of the available addresses. Yet the number of devices on the Internet
is much less than the 2^{32} address space. We have run out of class A and B addresses, and

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

a class C block is too small for most midsize organizations. One solution that has alleviated the problem is the idea of classless addressing.

> **Classful addressing, which is almost obsolete, is replaced with classless addressing.**

Classless Addressing

To overcome address depletion and give more organizations access to the Internet, **classless addressing** was designed and implemented. In this scheme, there are no classes, but the addresses are still granted in blocks.

Address Blocks

In classless addressing, when an entity, small or large, needs to be connected to the Internet, it is granted a block (range) of addresses. The size of the block (the number of addresses) varies based on the nature and size of the entity. For example, a household may be given only two addresses; a large organization may be given thousands of addresses. An ISP, as the Internet service provider, may be given thousands or hundreds of thousands based on the number of customers it may serve.

Restriction   To simplify the handling of addresses, the Internet authorities impose three restrictions on classless address blocks:

1. The addresses in a block must be contiguous, one after another.
2. The number of addresses in a block must be a power of 2 (1, 2, 4, 8, . . .).
3. The first address must be evenly divisible by the number of addresses.

Example 19.5

Figure 19.3 shows a block of addresses, in both binary and dotted-decimal notation, granted to a small business that needs 16 addresses.

Figure 19.3   *A block of 16 addresses granted to a small organization*

We can see that the restrictions are applied to this block. The addresses are contiguous. The number of addresses is a power of 2 ($16 = 2^4$), and the first address is divisible by 16. The first address, when converted to a decimal number, is 3,440,387,360, which when divided by 16 results in 215,024,210. In Appendix B, we show how to find the decimal value of an IP address.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Mask

A better way to define a block of addresses is to select any address in the block and the mask. As we discussed before, a mask is a 32-bit number in which the n leftmost bits are 1s and the $32 - n$ rightmost bits are 0s. However, in classless addressing the mask for a block can take any value from 0 to 32. It is very convenient to give just the value of n preceded by a slash (CIDR notation).

In IPv4 addressing, a block of addresses can be defined as
x.y.z.t /n
in which x.y.z.t defines one of the addresses and the /n defines the mask.

The address and the /n notation completely define the whole block (the first address, the last address, and the number of addresses).

First Address   The first address in the block can be found by setting the $32 - n$ rightmost bits in the binary notation of the address to 0s.

The first address in the block can be found by setting the rightmost $32 - n$ bits to 0s.

Example 19.6

A block of addresses is granted to a small organization. We know that one of the addresses is 205.16.37.39/28. What is the first address in the block?

Solution
The binary representation of the given address is 11001101 00010000 00100101 00100111. If we set $32 - 28$ rightmost bits to 0, we get 11001101 0001000 00100101 0010000 or 205.16.37.32. This is actually the block shown in Figure 19.3.

Last Address   The last address in the block can be found by setting the $32 - n$ rightmost bits in the binary notation of the address to 1s.

The last address in the block can be found by setting the rightmost $32 - n$ bits to 1s.

Example 19.7

Find the last address for the block in Example 19.6.

Solution
The binary representation of the given address is 11001101 00010000 00100101 00100111. If we set $32 - 28$ rightmost bits to 1, we get 11001101 00010000 00100101 00101111 or 205.16.37.47. This is actually the block shown in Figure 19.3.

Number of Addresses   The number of addresses in the block is the difference between the last and first address. It can easily be found using the formula 2^{32-n}.

The number of addresses in the block can be found by using the formula 2^{32-n}.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

SECTION 19.1   IPv4 ADDRESSES   **557**

Example 19.8

Find the number of addresses in Example 19.6.

Solution

The value of *n* is 28, which means that number of addresses is 2^{32-28} or 16.

Example 19.9

Another way to find the first address, the last address, and the number of addresses is to represent the mask as a 32-bit binary (or 8-digit hexadecimal) number. This is particularly useful when we are writing a program to find these pieces of information. In Example 19.5 the /28 can be represented as 11111111 11111111 11111111 11110000 (twenty-eight 1s and four 0s). Find

  a. The first address

  b. The last address

  c. The number of addresses

Solution

  a. The first address can be found by ANDing the given addresses with the mask. ANDing here is done bit by bit. The result of ANDing 2 bits is 1 if both bits are 1s; the result is 0 otherwise.

Address:	11001101 00010000 00100101 00100111
Mask:	**11111111 11111111 11111111 11110000**
First address:	11001101 00010000 00100101 00100000

  b. The last address can be found by ORing the given addresses with the complement of the mask. ORing here is done bit by bit. The result of ORing 2 bits is 0 if both bits are 0s; the result is 1 otherwise. The complement of a number is found by changing each 1 to 0 and each 0 to 1.

Address:	11001101 00010000 00100101 00100111
Mask complement:	**00000000 00000000 00000000 00001111**
Last address:	11001101 00010000 00100101 00101111

  c. The number of addresses can be found by complementing the mask, interpreting it as a decimal number, and adding 1 to it.

Mask complement:	**000000000 00000000 00000000 00001111**
Number of addresses:	15 + 1 = 16

Network Addresses

A very important concept in IP addressing is the **network address.** When an organization is given a block of addresses, the organization is free to allocate the addresses to the devices that need to be connected to the Internet. The first address in the class, however, is normally (not always) treated as a special address. The first address is called the network address and defines the organization network. It defines the organization itself to the rest of the world. In a later chapter we will see that the first address is the one that is used by routers to direct the message sent to the organization from the outside.

**Forouzan: Data
Communications and
Networking, Fourth Edition**

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

558   *CHAPTER 19   NETWORK LAYER: LOGICAL ADDRESSING*

Figure 19.4 shows an organization that is granted a 16-address block.

**Figure 19.4**   *A network configuration for the block 205.16.37.32/28*

The organization network is connected to the Internet via a router. The router has two addresses. One belongs to the granted block; the other belongs to the network that is at the other side of the router. We call the second address x.y.z.t/*n* because we do not know anything about the network it is connected to at the other side. All messages destined for addresses in the organization block (205.16.37.32 to 205.16.37.47) are sent, directly or indirectly, to x.y.z.t/*n*. We say directly or indirectly because we do not know the structure of the network to which the other side of the router is connected.

> **The first address in a block is normally not assigned to any device; it is used as the network address that represents the organization to the rest of the world.**

Hierarchy

IP addresses, like other addresses or identifiers we encounter these days, have levels of hierarchy. For example, a telephone network in North America has three levels of hierarchy. The leftmost three digits define the area code, the next three digits define the exchange, the last four digits define the connection of the local loop to the central office. Figure 19.5 shows the structure of a hierarchical telephone number.

Figure 19.5   *Hierarchy in a telephone network in North America*

Two-Level Hierarchy: No Subnetting

An IP address can define only two levels of hierarchy when not subnetted. The n leftmost bits of the address x.y.z.t/n define the network (organization network); the $32 - n$ rightmost bits define the particular host (computer or router) to the network. The two common terms are prefix and suffix. The part of the address that defines the network is called the **prefix;** the part that defines the host is called the **suffix.** Figure 19.6 shows the hierarchical structure of an IPv4 address.

Figure 19.6   *Two levels of hierarchy in an IPv4 address*

The prefix is common to all addresses in the network; the suffix changes from one device to another.

> Each address in the block can be considered as a two-level hierarchical structure:
> the leftmost n bits (prefix) define the network;
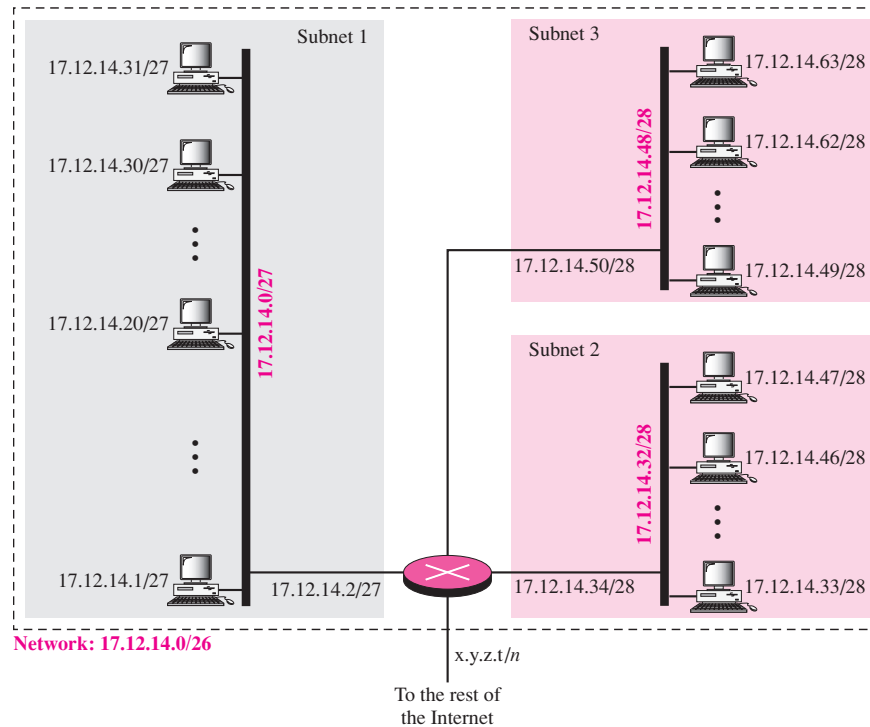> the rightmost $32 - n$ bits define the host.

Three-Levels of Hierarchy: Subnetting

An organization that is granted a large block of addresses may want to create clusters of networks (called subnets) and divide the addresses between the different subnets. The rest of the world still sees the organization as one entity; however, internally there are several subnets. All messages are sent to the router address that connects the organization to the rest of the Internet; the router routes the message to the appropriate subnets. The organization, however, needs to create small subblocks of addresses, each assigned to specific subnets. The organization has its own mask; each subnet must also have its own.

As an example, suppose an organization is given the block 17.12.40.0/26, which contains 64 addresses. The organization has three offices and needs to divide the addresses into three subblocks of 32, 16, and 16 addresses. We can find the new masks by using the following arguments:

1. Suppose the mask for the first subnet is n1, then 2^{32-n1} must be 32, which means that n1 = 27.
2. Suppose the mask for the second subnet is n2, then 2^{32-n2} must be 16, which means that n2 = 28.
3. Suppose the mask for the third subnet is n3, then 2^{32-n3} must be 16, which means that n3 = 28.

This means that we have the masks 27, 28, 28 with the organization mask being 26. Figure 19.7 shows one configuration for the above scenario.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Figure 19.7 *Configuration and addresses in a subnetted network*

Let us check to see if we can find the subnet addresses from one of the addresses in the subnet.

a. In subnet 1, the address 17.12.14.29/27 can give us the subnet address if we use the mask /27 because

```
Host:    00010001 00001100 00001110 00011101
Mask:    /27
Subnet:  00010001 00001100 00001110 00000000 ➡ (17.12.14.0)
```

b. In subnet 2, the address 17.12.14.45/28 can give us the subnet address if we use the mask /28 because

```
Host:    00010001 00001100 00001110 00101101
Mask:    /28
Subnet:  00010001 00001100 00001110 00100000 ➡ (17.12.14.32)
```

c. In subnet 3, the address 17.12.14.50/28 can give us the subnet address if we use the mask /28 because

```
Host:    00010001 00001100 00001110 00110010
Mask:    /28
Subnet:  00010001 00001100 00001110 00110000 ➡ (17.12.14.48)
```

Note that applying the mask of the network, /26, to any of the addresses gives us the network address 17.12.14.0/26. We leave this proof to the reader.

We can say that through subnetting, we have three levels of hierarchy. Note that in our example, the subnet prefix length can differ for the subnets as shown in Figure 19.8.

Figure 19.8   *Three-level hierarchy in an IPv4 address*

More Levels of Hierarchy

The structure of classless addressing does not restrict the number of hierarchical levels. An organization can divide the granted block of addresses into subblocks. Each subblock can in turn be divided into smaller subblocks. And so on. One example of this is seen in the ISPs. A national ISP can divide a granted large block into smaller blocks and assign each of them to a regional ISP. A regional ISP can divide the block received from the national ISP into smaller blocks and assign each one to a local ISP. A local ISP can divide the block received from the regional ISP into smaller blocks and assign each one to a different organization. Finally, an organization can divide the received block and make several subnets out of it.

Address Allocation

The next issue in classless addressing is address allocation. How are the blocks allocated? The ultimate responsibility of address allocation is given to a global authority called the *Internet Corporation for Assigned Names and Addresses* (ICANN). However, ICANN does not normally allocate addresses to individual organizations. It assigns a large block of addresses to an ISP. Each ISP, in turn, divides its assigned block into smaller subblocks and grants the subblocks to its customers. In other words, an ISP receives one large block to be distributed to its Internet users. This is called **address aggregation:** many blocks of addresses are aggregated in one block and granted to one ISP.

Example 19.10

An ISP is granted a block of addresses starting with 190.100.0.0/16 (65,536 addresses). The ISP needs to distribute these addresses to three groups of customers as follows:

   a. The first group has 64 customers; each needs 256 addresses.
   b. The second group has 128 customers; each needs 128 addresses.
   c. The third group has 128 customers; each needs 64 addresses.

Design the subblocks and find out how many addresses are still available after these allocations.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Solution

Figure 19.9 shows the situation.

Figure 19.9   *An example of address allocation and distribution by an ISP*
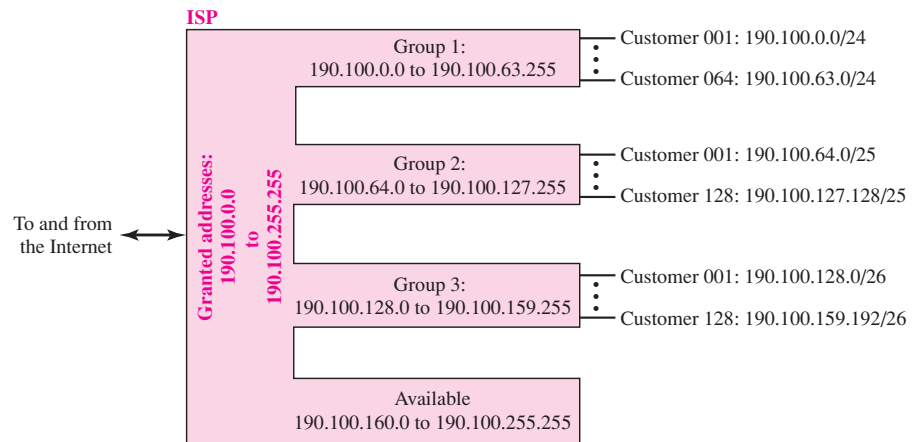
1. **Group 1**

   For this group, each customer needs 256 addresses. This means that 8 ($\log_2 256$) bits are needed to define each host. The prefix length is then $32 - 8 = 24$. The addresses are

   > *1st Customer:*   190.100.0.0/24   190.100.0.255/24
   > *2nd Customer:*   190.100.1.0/24   190.100.1.255/24
   >   . . .
   > *64th Customer:*   190.100.63.0/24   190.100.63.255/24
   > *Total = 64 × 256 = 16,384*

2. **Group 2**

   For this group, each customer needs 128 addresses. This means that 7 ($\log_2 128$) bits are needed to define each host. The prefix length is then $32 - 7 = 25$. The addresses are

   > *1st Customer:*   190.100.64.0/25   190.100.64.127/25
   > *2nd Customer:*   190.100.64.128/25   190.100.64.255/25
   >   . . .
   > *128th Customer:* 190.100.127.128/25   190.100.127.255/25
   > *Total = 128 × 128 = 16,384*

3. **Group 3**

   For this group, each customer needs 64 addresses. This means that 6 ($\log_2 64$) bits are needed to each host. The prefix length is then $32 - 6 = 26$. The addresses are

> 1st Customer: 190.100.128.0/26   190.100.128.63/26
> 2nd Customer: 190.100.128.64/26   190.100.128.127/26
>   . . .
> 128th Customer: 190.100.159.192/26   190.100.159.255/26
> Total = 128 × 64 = 8192

Number of granted addresses to the ISP: 65,536
Number of allocated addresses by the ISP: 40,960
Number of available addresses: 24,576

Network Address Translation (NAT)

The number of home users and small businesses that want to use the Internet is ever increasing. In the beginning, a user was connected to the Internet with a dial-up line, which means that she was connected for a specific period of time. An ISP with a block of addresses could dynamically assign an address to this user. An address was given to a user when it was needed. But the situation is different today. Home users and small businesses can be connected by an ADSL line or cable modem. In addition, many are not happy with one address; many have created small networks with several hosts and need an IP address for each host. With the shortage of addresses, this is a serious problem.

A quick solution to this problem is called **network address translation (NAT).** NAT enables a user to have a large set of addresses internally and one address, or a small set of addresses, externally. The traffic inside can use the large set; the traffic outside, the small set.

To separate the addresses used inside the home or business and the ones used for the Internet, the Internet authorities have reserved three sets of addresses as private addresses, shown in Table 19.3.

Table 19.3   *Addresses for private networks*

Range			Total
10.0.0.0	to	10.255.255.255	2^{24}
172.16.0.0	to	172.31.255.255	2^{20}
192.168.0.0	to	192.168.255.255	2^{16}

Any organization can use an address out of this set without permission from the Internet authorities. Everyone knows that these reserved addresses are for private networks. They are unique inside the organization, but they are not unique globally. No router will forward a packet that has one of these addresses as the destination address.

The site must have only one single connection to the global Internet through a router that runs the NAT software. Figure 19.10 shows a simple implementation of NAT.

As Figure 19.10 shows, the private network uses private addresses. The router that connects the network to the global address uses one private address and one global address. The private network is transparent to the rest of the Internet; the rest of the Internet sees only the NAT router with the address 200.24.5.8.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

**Figure 19.10**   *A NAT implementation*

Address Translation

All the outgoing packets go through the NAT router, which replaces the *source address* in the packet with the global NAT address. All incoming packets also pass through the NAT router, which replaces the *destination address* in the packet (the NAT router global address) with the appropriate private address. Figure 19.11 shows an example of address translation.

Figure 19.11   *Addresses in a NAT*

Translation Table

The reader may have noticed that translating the source addresses for outgoing packets is straightforward. But how does the NAT router know the destination address for a packet coming from the Internet? There may be tens or hundreds of private IP addresses, each belonging to one specific host. The problem is solved if the NAT router has a translation table.

Using One IP Address   In its simplest form, a translation table has only two columns: the private address and the external address (destination address of the packet). When the router translates the source address of the outgoing packet, it also makes note of the destination address—where the packet is going. When the response comes back from the destination, the router uses the source address of the packet (as the external address) to find the private address of the packet. Figure 19.12 shows the idea. Note that the addresses that are changed (translated) are shown in color.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Figure 19.12   *NAT address translation*

In this strategy, communication must always be initiated by the private network. The NAT mechanism described requires that the private network start the communication. As we will see, NAT is used mostly by ISPs which assign one single address to a customer. The customer, however, may be a member of a private network that has many private addresses. In this case, communication with the Internet is always initiated from the customer site, using a client program such as HTTP, TELNET, or FTP to access the corresponding server program. For example, when e-mail that originates from a non-customer site is received by the ISP e-mail server, the e-mail is stored in the mailbox of the customer until retrieved. A private network cannot run a server program for clients outside of its network if it is using NAT technology.

Using a Pool of IP Addresses   Since the NAT router has only one global address, only one private network host can access the same external host. To remove this restriction, the NAT router uses a pool of global addresses. For example, instead of using only one global address (200.24.5.8), the NAT router can use four addresses (200.24.5.8, 200.24.5.9, 200.24.5.10, and 200.24.5.11). In this case, four private network hosts can communicate with the same external host at the same time because each pair of addresses defines a connection. However, there are still some drawbacks. In this example, no more than four connections can be made to the same destination. Also, no private-network host can access two external server programs (e.g., HTTP and FTP) at the same time.

Using Both IP Addresses and Port Numbers   To allow a many-to-many relationship between private-network hosts and external server programs, we need more information in the translation table. For example, suppose two hosts with addresses 172.18.3.1 and 172.18.3.2 inside a private network need to access the HTTP server on external host

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

25.8.3.2. If the translation table has five columns, instead of two, that include the source and destination port numbers of the transport layer protocol, the ambiguity is eliminated. We discuss port numbers in Chapter 23. Table 19.4 shows an example of such a table.

Table 19.4   *Five-column translation table*

Private Address	Private Port	External Address	External Port	Transport Protocol
172.18.3.1	1400	25.8.3.2	80	TCP
172.18.3.2	1401	25.8.3.2	80	TCP
. . .	. . .	. . .	. . .	. . .

Note that when the response from HTTP comes back, the combination of source address (25.8.3.2) and destination port number (1400) defines the private network host to which the response should be directed. Note also that for this translation to work, the temporary port numbers (1400 and 1401) must be unique.

NAT and ISP

An ISP that serves dial-up customers can use NAT technology to conserve addresses. For example, suppose an ISP is granted 1000 addresses, but has 100,000 customers. Each of the customers is assigned a private network address. The ISP translates each of the 100,000 source addresses in outgoing packets to one of the 1000 global addresses; it translates the global destination address in incoming packets to the corresponding private address. Figure 19.13 shows this concept.

Figure 19.13   *An ISP and NAT*

19.2   IPv6 ADDRESSES

Despite all short-term solutions, such as classless addressing, Dynamic Host Configuration Protocol (DHCP), discussed in Chapter 21, and NAT, address depletion is still a long-term problem for the Internet. This and other problems in the IP protocol itself,

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

such as lack of accommodation for real-time audio and video transmission, and encryption and authentication of data for some applications, have been the motivation for IPv6. In this section, we compare the address structure of IPv6 to IPv4. In Chapter 20, we discuss both protocols.

Structure

An **IPv6 address** consists of 16 bytes (octets); it is 128 bits long.

> **An IPv6 address is 128 bits long.**

Hexadecimal Colon Notation

To make addresses more readable, IPv6 specifies **hexadecimal colon notation.** In this notation, 128 bits is divided into eight sections, each 2 bytes in length. Two bytes in hexadecimal notation requires four hexadecimal digits. Therefore, the address consists of 32 hexadecimal digits, with every four digits separated by a colon, as shown in Figure 19.14.

Figure 19.14   *IPv6 address in binary and hexadecimal colon notation*

Abbreviation

Although the IP address, even in hexadecimal format, is very long, many of the digits are zeros. In this case, we can abbreviate the address. The leading zeros of a section (four digits between two colons) can be omitted. Only the leading zeros can be dropped, not the trailing zeros (see Figure 19.15).

Figure 19.15   *Abbreviated IPv6 addresses*

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Using this form of **abbreviation,** 0074 can be written as 74, 000F as F, and 0000 as 0. Note that 3210 cannot be abbreviated. Further abbreviations are possible if there are consecutive sections consisting of zeros only. We can remove the zeros altogether and replace them with a double semicolon. Note that this type of abbreviation is allowed only once per address. If there are two runs of zero sections, only one of them can be abbreviated. Reexpansion of the abbreviated address is very simple: Align the unabbreviated portions and insert zeros to get the original expanded address.

Example 19.11

Expand the address 0:15::1:12:1213 to its original.

Solution

We first need to align the left side of the double colon to the left of the original pattern and the right side of the double colon to the right of the original pattern to find how many 0s we need to replace the double colon.

$$xxxx:xxxx:xxxx:xxxx:xxxx:xxxx:xxxx:xxxx$$
$$0:\ 15:\qquad\qquad\ :\quad 1:\ 12:1213$$

This means that the original address is

$$0000:0015:0000:0000:0000:0001:0012:1213$$

Address Space

IPv6 has a much larger address space; 2^{128} addresses are available. The designers of IPv6 divided the address into several categories. A few leftmost bits, called the *type prefix,* in each address define its category. The type prefix is variable in length, but it is designed such that no code is identical to the first part of any other code. In this way, there is no ambiguity; when an address is given, the type prefix can easily be determined. Table 19.5 shows the prefix for each type of address. The third column shows the fraction of each type of address relative to the whole address space.

Table 19.5   *Type prefixes for IPv6 addresses*

Type Prefix	Type	Fraction
0000 0000	Reserved	1/256
0000 0001	Unassigned	1/256
0000 001	ISO network addresses	1/128
0000 010	IPX (Novell) network addresses	1/128
0000 011	Unassigned	1/128
0000 1	Unassigned	1/32
0001	Reserved	1/16
001	Reserved	1/8
010	**Provider-based unicast addresses**	**1/8**

Table 19.5   *Type prefixes for IPv6 addresses (continued)*

Type Prefix	Type	Fraction
011	Unassigned	1/8
100	Geographic-based unicast addresses	1/8
101	Unassigned	1/8
110	Unassigned	1/8
1110	Unassigned	1/16
1111 0	Unassigned	1/32
1111 10	Unassigned	1/64
1111 110	Unassigned	1/128
1111 1110 0	Unassigned	1/512
1111 1110 10	Link local addresses	1/1024
1111 1110 11	Site local addresses	1/1024
1111 1111	Multicast addresses	1/256

Unicast Addresses

A **unicast address** defines a single computer. The packet sent to a unicast address must be delivered to that specific computer. IPv6 defines two types of unicast addresses: geographically based and provider-based. We discuss the second type here; the first type is left for future definition. The provider-based address is generally used by a normal host as a unicast address. The address format is shown in Figure 19.16.

Figure 19.16   *Prefixes for provider-based unicast address*
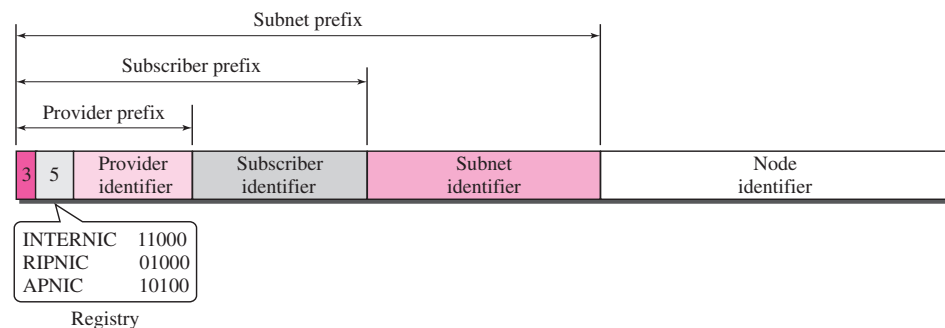
Fields for the provider-based address are as follows:

❏   **Type identifier.** This 3-bit field defines the address as a provider-based address.

❏   **Registry identifier.** This 5-bit field indicates the agency that has registered the address. Currently three registry centers have been defined. INTERNIC (code 11000) is the center for North America; RIPNIC (code 01000) is the center for European registration; and APNIC (code 10100) is for Asian and Pacific countries.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

❏ **Provider identifier.** This variable-length field identifies the provider for Internet access (such as an ISP). A 16-bit length is recommended for this field.

❏ **Subscriber identifier.** When an organization subscribes to the Internet through a provider, it is assigned a subscriber identification. A 24-bit length is recommended for this field.

❏ **Subnet identifier.** Each subscriber can have many different subnetworks, and each subnetwork can have an identifier. The subnet identifier defines a specific subnetwork under the territory of the subscriber. A 32-bit length is recommended for this field.

❏ **Node identifier.** The last field defines the identity of the node connected to a subnet. A length of 48 bits is recommended for this field to make it compatible with the 48-bit link (physical) address used by Ethernet. In the future, this link address will probably be the same as the node physical address.

Multicast Addresses

Multicast addresses are used to define a group of hosts instead of just one. A packet sent to a **multicast address** must be delivered to each member of the group. Figure 19.17 shows the format of a multicast address.

Figure 19.17   *Multicast address in IPv6*
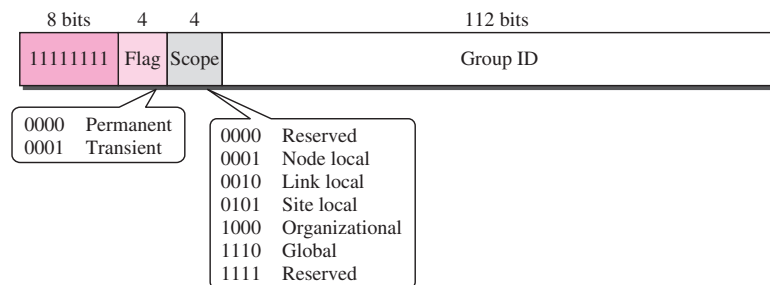
The second field is a flag that defines the group address as either permanent or transient. A permanent group address is defined by the Internet authorities and can be accessed at all times. A transient group address, on the other hand, is used only temporarily. Systems engaged in a teleconference, for example, can use a transient group address. The third field defines the scope of the group address. Many different scopes have been defined, as shown in Figure 19.17.

Anycast Addresses

IPv6 also defines anycast addresses. An **anycast address,** like a multicast address, also defines a group of nodes. However, a packet destined for an anycast address is delivered to only one of the members of the anycast group, the nearest one (the one with the shortest route). Although the definition of an anycast address is still debatable, one possible use is to assign an anycast address to all routers of an ISP that covers a large logical area in the Internet. The routers outside the ISP deliver a packet destined for the ISP to the nearest ISP router. No block is assigned for anycast addresses.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

Reserved Addresses

Another category in the address space is the **reserved address.** These addresses start with eight 0s (type prefix is 0000 0000). A few subcategories are defined in this category, as shown in Figure 19.18.

Figure 19.18   *Reserved addresses in IPv6*
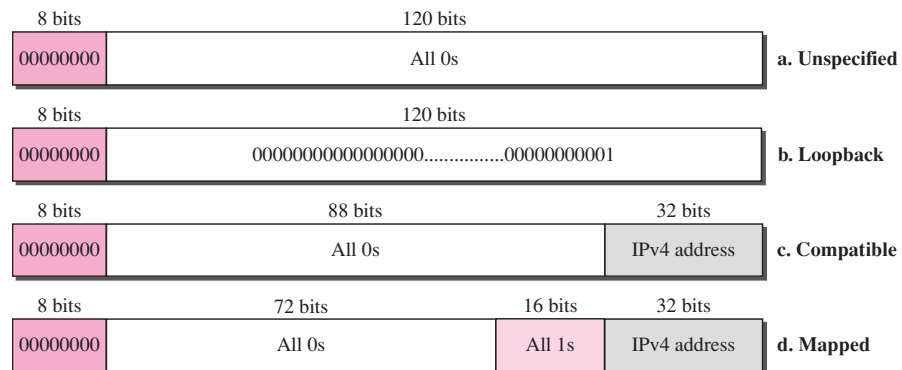
An **unspecified address** is used when a host does not know its own address and sends an inquiry to find its address. A **loopback address** is used by a host to test itself without going into the network. A **compatible address** is used during the transition from IPv4 to IPv6 (see Chapter 20). It is used when a computer using IPv6 wants to send a message to another computer using IPv6, but the message needs to pass through a part of the network that still operates in IPv4. A **mapped address** is also used during transition. However, it is used when a computer that has migrated to IPv6 wants to send a packet to a computer still using IPv4.

Local Addresses

These addresses are used when an organization wants to use IPv6 protocol without being connected to the global Internet. In other words, they provide addressing for private networks. Nobody outside the organization can send a message to the nodes using these addresses. Two types of addresses are defined for this purpose, as shown in Figure 19.19.

Figure 19.19   *Local addresses in IPv6*

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

A **link local address** is used in an isolated subnet; a **site local address** is used in an isolated site with several subnets.

19.3   RECOMMENDED READING

For more details about subjects discussed in this chapter, we recommend the following books and sites. The items in brackets [. . .] refer to the reference list at the end of the text.

Books

IPv4 addresses are discussed in Chapters 4 and 5 of [For06], Chapter 3 of [Ste94], Section 4.1 of [PD03], Chapter 18 of [Sta04], and Section 5.6 of [Tan03]. IPv6 addresses are discussed in Section 27.1 of [For06] and Chapter 8 of [Los04]. A good discussion of NAT can be found in [Dut01].

Sites

❏   www.ietf.org/rfc.html   Information about RFCs

RFCs

A discussion of IPv4 addresses can be found in most of the RFCs related to the IPv4 protocol:

> 760, 781, 791, 815, 1025, 1063, 1071, 1141, 1190, 1191, 1624, 2113

A discussion of IPv6 addresses can be found in most of the RFCs related to IPv6 protocol:

> 1365, 1550, 1678, 1680, 1682, 1683, 1686, 1688, 1726, 1752, 1826, 1883, 1884, 1886, 1887, 1955, 2080, 2373, 2452, 2463, 2465, 2466, 2472, 2492, 2545, 2590

A discussion of NAT can be found in

> 1361, 2663, 2694

19.4   KEY TERMS

address aggregation	class C address
address space	class D address
anycast address	class E address
binary notation	classful addressing
class A address	classless addressing
class B address	classless interdomain routing (CIDR)

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

SECTION 19.5 SUMMARY **573**

compatible address

dotted-decimal notation

default mask

hexadecimal colon notation

hostid

IP address

IPv4 address

IPv6 address

link local address

mapped address

mask

multicast address

netid

network address

network address translation
  (NAT)

prefix

reserved address

site local address

subnet

subnet mask

subnetting

suffix

supernet

supernet mask

supernetting

unicast address

unspecified address

19.5   SUMMARY

❏ At the network layer, a global identification system that uniquely identifies every host and router is necessary for delivery of a packet from host to host.

❏ An IPv4 address is 32 bits long and uniquely and universally defines a host or router on the Internet.

❏ In classful addressing, the portion of the IP address that identifies the network is called the netid.

❏ In classful addressing, the portion of the IP address that identifies the host or router on the network is called the hostid.

❏ An IP address defines a device's connection to a network.

❏ There are five classes in IPv4 addresses. Classes A, B, and C differ in the number of hosts allowed per network. Class D is for multicasting and Class E is reserved.

❏ The class of an address is easily determined by examination of the first byte.

❏ Addresses in classes A, B, or C are mostly used for unicast communication.

❏ Addresses in class D are used for multicast communication.

❏ Subnetting divides one large network into several smaller ones, adding an intermediate level of hierarchy in IP addressing.

❏ Supernetting combines several networks into one large one.

❏ In classless addressing, we can divide the address space into variable-length blocks.

❏ There are three restrictions in classless addressing:

    a. The number of addresses needs to be a power of 2.

    b. The mask needs to be included in the address to define the block.

    c. The starting address must be divisible by the number of addresses in the block.

❏ The mask in classless addressing is expressed as the prefix length (/*n*) in CIDR notation.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

❑ To find the first address in a block, we set the rightmost $32 - n$ bits to 0.

❑ To find the number of addresses in the block, we calculate 2^{32-n}, where n is the prefix length.

❑ To find the last address in the block, we set the rightmost $32 - n$ bits to 0.

❑ Subnetting increases the value of n.

❑ The global authority for address allocation is ICANN. ICANN normally grants large blocks of addresses to ISPs, which in turn grant small subblocks to individual customers.

❑ IPv6 addresses use hexadecimal colon notation with abbreviation methods available.

❑ There are three types of addresses in IPv6: unicast, anycast, and multicast.

❑ In an IPv6 address, the variable type prefix field defines the address type or purpose.

19.6   PRACTICE SET

Review Questions

1. What is the number of bits in an IPv4 address? What is the number of bits in an IPv6 address?

2. What is dotted decimal notation in IPv4 addressing? What is the number of bytes in an IPv4 address represented in dotted decimal notation? What is hexadecimal notation in IPv6 addressing? What is the number of digits in an IPv6 address represented in hexadecimal notation?

3. What are the differences between classful addressing and classless addressing in IPv4?

4. List the classes in classful addressing and define the application of each class (unicast, multicast, broadcast, or reserve).

5. Explain why most of the addresses in class A are wasted. Explain why a medium-size or large-size corporation does not want a block of class C addresses.

6. What is a mask in IPv4 addressing? What is a default mask in IPv4 addressing?

7. What is the network address in a block of addresses? How can we find the network address if one of the addresses in a block is given?

8. Briefly define subnetting and supernetting. How do the subnet mask and supernet mask differ from a default mask in classful addressing?

9. How can we distinguish a multicast address in IPv4 addressing? How can we do so in IPv6 addressing?

10. What is NAT? How can NAT help in address depletion?

Exercises

11. What is the address space in each of the following systems?

    a. A system with 8-bit addresses

    b. A system with 16-bit addresses

    c. A system with 64-bit addresses

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

12. An address space has a total of 1024 addresses. How many bits are needed to represent an address?

13. An address space uses the three symbols 0, 1, and 2 to represent addresses. If each address is made of 10 symbols, how many addresses are available in this system?

14. Change the following IP addresses from dotted-decimal notation to binary notation.
    a. 114.34.2.8
    b. 129.14.6.8
    c. 208.34.54.12
    d. 238.34.2.1

15. Change the following IP addresses from binary notation to dotted-decimal notation.
    a. 01111111 11110000 01100111 01111101
    b. 10101111 11000000 11111000 00011101
    c. 11011111 10110000 00011111 01011101
    d. 11101111 11110111 11000111 00011101

16. Find the class of the following IP addresses.
    a. 208.34.54.12
    b. 238.34.2.1
    c. 114.34.2.8
    d. 129.14.6.8

17. Find the class of the following IP addresses.
    a. 11110111 11110011 10000111 11011101
    b. 10101111 11000000 11110000 00011101
    c. 11011111 10110000 00011111 01011101
    d. 11101111 11110111 11000111 00011101

18. Find the netid and the hostid of the following IP addresses.
    a. 114.34.2.8
    b. 132.56.8.6
    c. 208.34.54.12

19. In a block of addresses, we know the IP address of one host is 25.34.12.56/16. What are the first address (network address) and the last address (limited broadcast address) in this block?

20. In a block of addresses, we know the IP address of one host is 182.44.82.16/26. What are the first address (network address) and the last address in this block?

21. An organization is granted the block 16.0.0.0/8. The administrator wants to create 500 fixed-length subnets.
    a. Find the subnet mask.
    b. Find the number of addresses in each subnet.
    c. Find the first and last addresses in subnet 1.
    d. Find the first and last addresses in subnet 500.

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

22. An organization is granted the block 130.56.0.0/16. The administrator wants to create 1024 subnets.
   a. Find the subnet mask.
   b. Find the number of addresses in each subnet.
   c. Find the first and last addresses in subnet 1.
   d. Find the first and last addresses in subnet 1024.

23. An organization is granted the block 211.17.180.0/24. The administrator wants to create 32 subnets.
   a. Find the subnet mask.
   b. Find the number of addresses in each subnet.
   c. Find the first and last addresses in subnet 1.
   d. Find the first and last addresses in subnet 32.

24. Write the following masks in slash notation (/*n*).
   a. 255.255.255.0
   b. 255.0.0.0
   c. 255.255.224.0
   d. 255.255.240.0

25. Find the range of addresses in the following blocks.
   a. 123.56.77.32/29
   b. 200.17.21.128/27
   c. 17.34.16.0/23
   d. 180.34.64.64/30

26. An ISP is granted a block of addresses starting with 150.80.0.0/16. The ISP wants to distribute these blocks to 2600 customers as follows.
   a. The first group has 200 medium-size businesses; each needs 128 addresses.
   b. The second group has 400 small businesses; each needs 16 addresses.
   c. The third group has 2000 households; each needs 4 addresses.

   Design the subblocks and give the slash notation for each subblock. Find out how many addresses are still available after these allocations.

27. An ISP is granted a block of addresses starting with 120.60.4.0/22. The ISP wants to distribute these blocks to 100 organizations with each organization receiving just eight addresses. Design the subblocks and give the slash notation for each subblock. Find out how many addresses are still available after these allocations.

28. An ISP has a block of 1024 addresses. It needs to divide the addresses among 1024 customers. Does it need subnetting? Explain your answer.

29. Show the shortest form of the following addresses.
   a. 2340:1ABC:119A:A000:0000:0000:0000:0000
   b. 0000:00AA:0000:0000:0000:0000:119A:A231
   c. 2340:0000:0000:0000:0000:119A:A001:0000
   d. 0000:0000:0000:2340:0000:0000:0000:0000

Forouzan: Data
Communications and
Networking, Fourth Edition

IV. Network Layer

19. Network Layer: Logical
Addressing

© The McGraw–Hill
Companies, 2007

30. Show the original (unabbreviated) form of the following addresses.
    a. 0::0
    b. 0:AA::0
    c. 0:1234::3
    d. 123::1:2
31. What is the type of each of the following addresses?
    a. FE80::12
    b. FEC0::24A2
    c. FF02::0
    d. 0::01
32. What is the type of each of the following addresses?
    a. 0::0
    b. 0::FFFF:0:0
    c. 582F:1234::2222
    d. 4821::14:22
    e. 54EF::A234:2
33. Show the provider prefix (in hexadecimal colon notation) of an address assigned to a subscriber if it is registered in the United States with ABC1 as the provider identification.
34. Show in hexadecimal colon notation the IPv6 address
    a. Compatible to the IPv4 address 129.6.12.34
    b. Mapped to the IPv4 address 129.6.12.34
35. Show in hexadecimal colon notation
    a. The link local address in which the node identifier is 0::123/48
    b. The site local address in which the node identifier is 0::123/48
36. Show in hexadecimal colon notation the permanent multicast address used in a link local scope.
37. A host has the address 581E:1456:2314:ABCD::1211. If the node identification is 48 bits, find the address of the subnet to which the host is attached.
38. A site with 200 subnets has the class B address of 132.45.0.0. The site recently migrated to IPv6 with the subscriber prefix 581E:1456:2314::ABCD/80. Design the subnets and define the subnet addresses, using a subnet identifier of 32 bits.

Research Activities

39. Find the block of addresses assigned to your organization or institution.
40. If you are using an ISP to connect from your home to the Internet, find the name of the ISP and the block of addresses assigned to it.
41. Some people argue that we can consider the whole address space as one single block in which each range of addresses is a subblock to this single block. Elaborate on this idea. What happens to subnetting if we accept this concept?
42. Is your school or organization using a classful address? If so, find out the class of the address.